

[kernel](#) / [pub](#) / [scm](#) / [fs](#) / [xfs](#) / [xfs-documentation](#) / [master](#) / [.](#) / [design](#) / [XFS_FileSystem_Structure](#) / **allocation_groups.asciidoc**

blob: c0ba16a82d489fedc10bfafd90c26409938b7730 [[file](#)] [[log](#)] [[blame](#)]

```
1  [[Allocation_Groups]]
2  = Allocation Groups
3
4  As mentioned earlier, XFS filesystems are divided into a number of equally
5  sized chunks called Allocation Groups. Each AG can almost be thought of as an
6  individual filesystem that maintains its own space usage. Each AG can be up to
7  one terabyte in size (512 bytes × 231), regardless of the underlying device's
8  sector size.
9
10 Each AG has the following characteristics:
11
12     * A super block describing overall filesystem info
13     * Free space management
14     * Inode allocation and tracking
15     * Reverse block-mapping index (optional)
16     * Data block reference count index (optional)
17
18 Having multiple AGs allows XFS to handle most operations in parallel without
19 degrading performance as the number of concurrent accesses increases.
20
21 The only global information maintained by the first AG (primary) is free space
22 across the filesystem and total inode counts. If the
23 +XFS_SB_VERSION2_LAZYSBCOUNTBIT+ flag is set in the superblock, these are only
24 updated on-disk when the filesystem is cleanly unmounted (umount or shutdown).
25
26 Immediately after a +mkfs.xfs+, the primary AG has the following disk layout;
27 the subsequent AGs do not have any inodes allocated:
28
29 .Allocation group layout
30 image::images/6.png[]
31
32 Each of these structures are expanded upon in the following sections.
33
34 [[Superblocks]]
35 == Superblocks
36
37 Each AG starts with a superblock. The first one, in AG 0, is the primary
38 superblock which stores aggregate AG information. Secondary superblocks are
39 only used by xfs_repair when the primary superblock has been corrupted. A
40 superblock is one sector in length.
```

The superblock is defined by the following structure. The description of each field follows.

[source, c]

struct xfs_sb

{

__uint32_t	sb_magicnum;
__uint32_t	sb_blocksize;
xfs_rfsblock_t	sb_dblocks;
xfs_rfsblock_t	sb_rblocks;
xfs_rtbblock_t	sb_rextents;
uuid_t	sb_uuid;
xfs_fsblock_t	sb_logstart;
xfs_ino_t	sb_rootino;
xfs_ino_t	sb_rbmino;
xfs_ino_t	sb_rsumino;
xfs_agblock_t	sb_rextsize;
xfs_agblock_t	sb_agblocks;
xfs_agnumber_t	sb_agcount;
xfs_extlen_t	sb_rbmblocks;
xfs_extlen_t	sb_logblocks;
__uint16_t	sb_versionnum;
__uint16_t	sb_sectsize;
__uint16_t	sb_inodesize;
__uint16_t	sb_inopblock;
char	sb_fname[12];
__uint8_t	sb_blocklog;
__uint8_t	sb_sectlog;
__uint8_t	sb_inodelog;
__uint8_t	sb_inopblog;
__uint8_t	sb_agblklog;
__uint8_t	sb_rextslog;
__uint8_t	sb_inprogress;
__uint8_t	sb_imax_pct;
__uint64_t	sb_icount;
__uint64_t	sb_ifree;
__uint64_t	sb_fdblocks;
__uint64_t	sb_frextents;
xfs_ino_t	sb_uquotino;
xfs_ino_t	sb_gquotino;
__uint16_t	sb_qflags;
__uint8_t	sb_flags;
__uint8_t	sb_shared_vn;
xfs_extlen_t	sb_inoalignmt;
__uint32_t	sb_unit;
__uint32_t	sb_width;
__uint8_t	sb_dirblklog;

```

90         __uint8_t          sb_logsectlog;
91         __uint16_t         sb_logsectsize;
92         __uint32_t         sb_logsunit;
93         __uint32_t         sb_features2;
94         __uint32_t         sb_bad_features2;
95
96         /* version 5 superblock fields start here */
97         __uint32_t         sb_features_compat;
98         __uint32_t         sb_features_ro_compat;
99         __uint32_t         sb_features_incompat;
100        __uint32_t         sb_features_log_incompat;
101
102        __uint32_t         sb_crc;
103        xfs_extlen_t       sb_spino_align;
104
105        xfs_ino_t          sb_pquotino;
106        xfs_lsn_t          sb_lsn;
107        uuid_t             sb_meta_uuid;
108        xfs_ino_t          sb_rrmapino;
109    };
110    ----
111    *sb_magicnum*::
112    Identifies the filesystem. Its value is +XFS_SB_MAGIC+ ``XFSB'' (0x58465342).
113
114    *sb_blocksize*::
115    The size of a basic unit of space allocation in bytes. Typically, this is 4096
116    (4KB) but can range from 512 to 65536 bytes.
117
118    *sb_dblocks*::
119    Total number of blocks available for data and metadata on the filesystem.
120
121    *sb_rblocks*::
122    Number blocks in the real-time disk device. Refer to
123    xref:Real-time_Devices[real-time sub-volumes] for more information.
124
125    *sb_rextents*::
126    Number of extents on the real-time device.
127
128    *sb_uuid*::
129    UUID (Universally Unique ID) for the filesystem. Filesystems can be mounted by
130    the UUID instead of device name.
131
132    *sb_logstart*::
133    First block number for the journaling log if the log is internal (ie. not on a
134    separate disk device). For an external log device, this will be zero (the log
135    will also start on the first block on the log device). The identity of the log
136    devices is not recorded in the filesystem, but the UUIDs of the filesystem and
137    the log device are compared to prevent corruption.
138

```

```

139  *sb_rootino*::
140  Root inode number for the filesystem. Normally, the root inode is at the
141  start of the first possible inode chunk in AG 0. This is 128 when using a 4KB
142  block size.
143
144  *sb_rbmino*::
145  Bitmap inode for real-time extents.
146
147  *sb_rsumino*::
148  Summary inode for real-time bitmap.
149
150  *sb_rextsize*::
151  Realtime extent size in blocks.
152
153  *sb_agblocks*::
154  Size of each AG in blocks. For the actual size of the last AG, refer to the
155  xref:AG_Free_Space_Management[free space] +agf_length+ value.
156
157  *sb_agcount*::
158  Number of AGs in the filesystem.
159
160  *sb_rmblocks*::
161  Number of real-time bitmap blocks.
162
163  *sb_logblocks*::
164  Number of blocks for the journaling log.
165
166  *sb_versionnum*::
167  Filesystem version number. This is a bitmask specifying the features enabled
168  when creating the filesystem. Any disk checking tools or drivers that do not
169  recognize any set bits must not operate upon the filesystem. Most of the flags
170  indicate features introduced over time. If the value of the lower nibble is >=
171  4, the higher bits indicate feature flags as follows:
172
173  .Version 4 Superblock version flags
174  [options="header"]
175  |====
176  | Flag                                | Description
177  | +XFS_SB_VERSION_ATTRBIT+           |
178  Set if any inode have extended attributes. If this bit is set; the
179  +XFS_SB_VERSION2_ATTR2BIT+ is not set; and the +attr2+ mount flag is not
180  specified, the +di_forkoff+ inode field will not be dynamically adjusted.
181  See the section about xref:Extended_Attribute_Versions[extended attribute
182  versions] for more information.
183
184  | +XFS_SB_VERSION_NLINKBIT+          | Set if any inodes use 32-bit di_nlink values.
185  | +XFS_SB_VERSION_QUOTABIT+          |
186  Quotas are enabled on the filesystem. This
187  also brings in the various quota fields in the superblock.

```

```

188
189 | +XFS_SB_VERSION_ALIGNBIT+ | Set if sb_inoalignmt is used.
190 | +XFS_SB_VERSION_DALIGNBIT+ | Set if sb_unit and sb_width are used.
191 | +XFS_SB_VERSION_SHAREDBIT+ | Set if sb_shared_vn is used.
192 | +XFS_SB_VERSION_LOGV2BIT+ | Version 2 journaling logs are used.
193 | +XFS_SB_VERSION_SECTORBIT+ | Set if sb_sectsize is not 512.
194 | +XFS_SB_VERSION_EXTFLGBIT+ | Unwritten extents are used. This is always set.
195 | +XFS_SB_VERSION_DIRV2BIT+ |
196 Version 2 directories are used. This is always set.
197
198 | +XFS_SB_VERSION_MOREBITSBIT+ |
199 Set if the sb_features2 field in the superblock contains more flags.
200 |====
201
202 If the lower nibble of this value is 5, then this is a v5 filesystem; the
203 +XFS_SB_VERSION2_CRCBIT+ feature must be set in +sb_features2+.
204
205 *sb_sectsize*::
206 Specifies the underlying disk sector size in bytes. Typically this is 512 or
207 4096 bytes. This determines the minimum I/O alignment, especially for direct I/O.
208
209 *sb_inodesize*::
210 Size of the inode in bytes. The default is 256 (2 inodes per standard sector)
211 but can be made as large as 2048 bytes when creating the filesystem. On a v5
212 filesystem, the default and minimum inode size are both 512 bytes.
213
214 *sb_inopblock*::
215 Number of inodes per block. This is equivalent to +sb_blocksize / sb_inodesize+.
216
217 *sb_fname[12]*::
218 Name for the filesystem. This value can be used in the mount command.
219
220 *sb_blocklog*::
221 log~2~ value of +sb_blocksize+. In other terms, +sb_blocksize = 2^sb_blocklog^+.
222
223 *sb_sectlog*::
224 log~2~ value of +sb_sectsize+.
225
226 *sb_inodelog*::
227 log~2~ value of +sb_inodesize+.
228
229 *sb_inopblog*::
230 log~2~ value of +sb_inopblock+.
231
232 *sb_agblklog*::
233 log~2~ value of +sb_agblocks+ (rounded up). This value is used to generate inode
234 numbers and absolute block numbers defined in extent maps.
235
236 *sb_rextslog*::

```

```

237 log~2~ value of +sb_rextents+.
238
239 *sb_inprogress*::
240 Flag specifying that the filesystem is being created.
241
242 *sb_imax_pct*::
243 Maximum percentage of filesystem space that can be used for inodes. The default
244 value is 5%.
245
246 *sb_icount*::
247 Global count for number inodes allocated on the filesystem. This is only
248 maintained in the first superblock.
249
250 *sb_ifree*::
251 Global count of free inodes on the filesystem. This is only maintained in the
252 first superblock.
253
254 *sb_fdblocks*::
255 Global count of free data blocks on the filesystem. This is only maintained in
256 the first superblock.
257
258 *sb_frextents*::
259 Global count of free real-time extents on the filesystem. This is only
260 maintained in the first superblock.
261
262 *sb_uquotino*::
263 Inode for user quotas. This and the following two quota fields only apply if
264 +XFS_SB_VERSION_QUOTABIT+ flag is set in +sb_versionnum+. Refer to
265 xref:Quota_Inodes[quota inodes] for more information.
266
267 *sb_gquotino*::
268 Inode for group or project quotas. Group and project quotas cannot be used at
269 the same time on v4 filesystems. On a v5 filesystem, this inode always stores
270 group quota information.
271
272 *sb_qflags*::
273 Quota flags. It can be a combination of the following flags:
274
275 .Superblock quota flags
276 [options="header"]
277 |=====
278 | Flag                                | Description
279 | +XFS_UQUOTA_ACCT+                   | User quota accounting is enabled.
280 | +XFS_UQUOTA_ENFD+                   | User quotas are enforced.
281 | +XFS_UQUOTA_CHKD+                   | User quotas have been checked.
282 | +XFS_PQUOTA_ACCT+                   | Project quota accounting is enabled.
283 | +XFS_OQUOTA_ENFD+                   | Other (group/project) quotas are enforced.
284 | +XFS_OQUOTA_CHKD+                   | Other (group/project) quotas have been checked.
285 | +XFS_GQUOTA_ACCT+                   | Group quota accounting is enabled.

```

```

286 | +XFS_GQUOTA_ENFD+ | Group quotas are enforced.
287 | +XFS_GQUOTA_CHKD+ | Group quotas have been checked.
288 | +XFS_PQUOTA_ENFD+ | Project quotas are enforced.
289 | +XFS_PQUOTA_CHKD+ | Project quotas have been checked.
290 |====
291
292 *sb_flags*::
293 Miscellaneous flags.
294
295 .Superblock flags
296 [options="header"]
297 |====
298 | Flag | Description
299 | +XFS_SBF_READONLY+ | Only read-only mounts allowed.
300 |====
301
302 *sb_shared_vn*::
303 Reserved and must be zero (``vn'' stands for version number).
304
305 *sb_inoalignmt*::
306 Inode chunk alignment in fsblocks. Prior to v5, the default value provided for
307 inode chunks to have an 8KiB alignment. Starting with v5, the default value
308 scales with the multiple of the inode size over 256 bytes. Concretely, this
309 means an alignment of 16KiB for 512-byte inodes, 32KiB for 1024-byte inodes,
310 etc. If sparse inodes are enabled, the +ir_startino+ field of each inode
311 B+tree record must be aligned to this block granularity, even if the inode
312 given by +ir_startino+ itself is sparse.
313
314 *sb_unit*::
315 Underlying stripe or raid unit in blocks.
316
317 *sb_width*::
318 Underlying stripe or raid width in blocks.
319
320 *sb_dirblklog*::
321 log~2~ multiplier that determines the granularity of directory block allocations
322 in fsblocks.
323
324 *sb_logsectlog*::
325 log~2~ value of the log subvolume's sector size. This is only used if the
326 journaling log is on a separate disk device (i.e. not internal).
327
328 *sb_logsectsize*::
329 The log's sector size in bytes if the filesystem uses an external log device.
330
331 *sb_logsunit*::
332 The log device's stripe or raid unit size. This only applies to version 2 logs
333 +XFS_SB_VERSION_LOGV2BIT+ is set in +sb_versionnum+.
334

```

```

335  *sb_features2*::
336  Additional version flags if +XFS_SB_VERSION_MOREBITSBIT+ is set in
337  +sb_versionnum+. The currently defined additional features include:
338
339  .Extended Version 4 Superblock flags
340  [options="header"]
341  |====
342  | Flag                                | Description
343  | +XFS_SB_VERSION2_LAZYSBCOUNTBIT+ |
344  Lazy global counters. Making a filesystem with this bit set can improve
345  performance. The global free space and inode counts are only updated in the
346  primary superblock when the filesystem is cleanly unmounted.
347
348  | +XFS_SB_VERSION2_ATTR2BIT+      |
349  Extended attributes version 2. Making a filesystem with this optimises the
350  inode layout of extended attributes. If this bit is set and the +noattr2+
351  mount flag is not specified, the +di_forkoff+ inode field will be dynamically
352  adjusted. See the section about xref:Extended_Attribute_Versions[extended
353  attribute versions] for more information.
354
355  | +XFS_SB_VERSION2_PARENTBIT+    |
356  Parent pointers. All inodes must have an extended attribute that points back to
357  its parent inode. The primary purpose for this information is in backup systems.
358
359  | +XFS_SB_VERSION2_PROJID32BIT+ |
360  32-bit Project ID. Inodes can be associated with a project ID number, which
361  can be used to enforce disk space usage quotas for a particular group of
362  directories. This flag indicates that project IDs can be 32 bits in size.
363
364  | +XFS_SB_VERSION2_CRCBIT+       |
365  Metadata checksumming. All metadata blocks have an extended header containing
366  the block checksum, a copy of the metadata UUID, the log sequence number of the
367  last update to prevent stale replays, and a back pointer to the owner of the
368  block. This feature must be and can only be set if the lowest nibble of
369  +sb_versionnum+ is set to 5.
370
371  | +XFS_SB_VERSION2_FTYPE+        |
372  Directory file type. Each directory entry records the type of the inode to
373  which the entry points. This speeds up directory iteration by removing the
374  need to load every inode into memory.
375  |====
376
377  *sb_bad_features2*::
378  This field mirrors +sb_features2+, due to past 64-bit alignment errors.
379
380  *sb_features_compat*::
381  Read-write compatible feature flags. The kernel can still read and write this
382  FS even if it doesn't understand the flag. Currently, there are no valid
383  flags.

```



```

384
385 *sb_features_ro_compat*::
386 Read-only compatible feature flags. The kernel can still read this FS even if
387 it doesn't understand the flag.
388
389 .Extended Version 5 Superblock Read-Only compatibility flags
390 [options="header"]
391 |====
392 | Flag | Description
393 | +XFS_SB_FEAT_RO_COMPAT_FINOBT+ |
394 Free inode B+tree. Each allocation group contains a B+tree to track inode chunks
395 containing free inodes. This is a performance optimization to reduce the time
396 required to allocate inodes.
397
398 | +XFS_SB_FEAT_RO_COMPAT_RMAPBT+ |
399 Reverse mapping B+tree. Each allocation group contains a B+tree containing
400 records mapping AG blocks to their owners. See the section about
401 xref:Reconstruction[reconstruction] for more details.
402
403 | +XFS_SB_FEAT_RO_COMPAT_REFLINK+ |
404 Reference count B+tree. Each allocation group contains a B+tree to track the
405 reference counts of AG blocks. This enables files to share data blocks safely.
406 See the section about xref:Reflink_Deduplication[reflink and deduplication] for
407 more details.
408
409 | +XFS_SB_FEAT_RO_COMPAT_INOBT CNT+ |
410 Inode B+tree block counters. Each allocation group's inode (AGI) header
411 tracks the number of blocks in each of the inode B+trees. This allows us
412 to have a slightly higher level of redundancy over the shape of the inode
413 btrees, and decreases the amount of time to compute the metadata B+tree
414 preallocations at mount time.
415
416 |====
417
418 *sb_features_incompat*::
419 Read-write incompatible feature flags. The kernel cannot read or write this
420 FS if it doesn't understand the flag.
421
422 .Extended Version 5 Superblock Read-Write incompatibility flags
423 [options="header"]
424 |====
425 | Flag | Description
426 | +XFS_SB_FEAT_INCOMPAT_FTYPE+ |
427 Directory file type. Each directory entry tracks the type of the inode to
428 which the entry points. This is a performance optimization to remove the need
429 to load every inode into memory to iterate a directory.
430
431 | +XFS_SB_FEAT_INCOMPAT_SPINODES+ |
432 Sparse inodes. This feature relaxes the requirement to allocate inodes in

```

```

433 chunks of 64. When the free space is heavily fragmented, there might exist
434 plenty of free space but not enough contiguous free space to allocate a new
435 inode chunk. With this feature, the user can continue to create files until
436 all free space is exhausted.
437
438 Unused space in the inode B+tree records are used to track which parts of the
439 inode chunk are not inodes.
440
441 See the chapter on xref:Sparse_Inodes[Sparse Inodes] for more information.
442
443 | +XFS_SB_FEAT_INCOMPAT_META_UUID+ |
444 Metadata UUID. The UUID stamped into each metadata block must match the value
445 in +sb_meta_uuid+. This enables the administrator to change +sb_uuid+ at will
446 without having to rewrite the entire filesystem.
447
448 | +XFS_SB_FEAT_INCOMPAT_BIGTIME+ |
449 Large timestamps. Inode timestamps and quota expiration timers are extended to
450 support times through the year 2486. See the section on
451 xref:Timestamps[timestamps] for more information.
452
453 | +XFS_SB_FEAT_INCOMPAT_NEEDSREPAIR+ |
454 The filesystem is not in operable condition, and must be run through
455 xfs_repair before it can be mounted.
456
457 | +XFS_SB_FEAT_INCOMPAT_NREXT64+ |
458 Large file fork extent counts. This greatly expands the maximum number of
459 space mappings allowed in data and extended attribute file forks.
460
461 |====
462
463 *sb_features_log_incompat*::
464 Read-write incompatible feature flags for the log. The kernel cannot recover
465 the FS log if it doesn't understand the flag.
466
467 .Extended Version 5 Superblock Log incompatibility flags
468 [options="header"]
469 |====
470 | Flag | Description
471 | +XFS_SB_FEAT_INCOMPAT_LOG_XATTRS+ |
472 Extended attribute updates have been committed to the ondisk log.
473
474 |====
475
476 *sb_crc*::
477 Superblock checksum.
478
479 *sb_spino_align*::
480 Sparse inode alignment, in fsblocks. Each chunk of inodes referenced by a
481 sparse inode B+tree record must be aligned to this block granularity.

```

```

482
483 *sb_pquotino*::
484 Project quota inode.
485
486 *sb_lsn*::
487 Log sequence number of the last superblock update.
488
489 *sb_meta_uuid*::
490 If the +XFS_SB_FEAT_INCOMPAT_META_UUID+ feature is set, then the UUID field in
491 all metadata blocks must match this UUID. If not, the block header UUID field
492 must match +sb_uuid+.
493
494 *sb_rrmapino*::
495 If the +XFS_SB_FEAT_RO_COMPAT_RMAPBT+ feature is set and a real-time
496 device is present (+sb_rblocks+ > 0), this field points to an inode
497 that contains the root to the
498 xref:Real_time_Reverse_Mapping_Btree[Real-Time Reverse Mapping B+tree].
499 This field is zero otherwise.
500
501 === xfs_db Superblock Example
502
503 A filesystem is made on a single disk with the following command:
504
505 ----
506 # mkfs.xfs -i attr=2 -n size=16384 -f /dev/sda7
507 meta-data=/dev/sda7          isize=256    agcount=16, agsize=3923122 blks
508              =               sectsz=512    attr=2
509 data      =                  bsize=4096    blocks=62769952, imaxpct=25
510              =               sunit=0      swidth=0 blks, unwritten=1
511 naming    =version 2         bsize=16384
512 log       =internal log      bsize=4096    blocks=30649, version=1
513              =               sectsz=512    sunit=0 blks
514 realtime  =none              extsz=65536   blocks=0, rtextents=0
515 ----
516
517 And in xfs_db, inspecting the superblock:
518
519 ----
520 xfs_db> sb
521 xfs_db> p
522 magicnum = 0x58465342
523 blocksize = 4096
524 dblocks = 62769952
525 rblocks = 0
526 rtextents = 0
527 uuid = 32b24036-6931-45b4-b68c-cd5e7d9a1ca5
528 logstart = 33554436
529 rootino = 128
530 rbmino = 129

```

```
531     rsumino = 130
532     rextsize = 16
533     agblocks = 3923122
534     agcount = 16
535     rbmblocks = 0
536     logblocks = 30649
537     versionnum = 0xb084
538     sectsize = 512
539     inodesize = 256
540     inopblock = 16
541     fname = "\000\000\000\000\000\000\000\000\000\000\000\000"
542     blocklog = 12
543     sectlog = 9
544     inodelog = 8
545     inopblog = 4
546     agblklog = 22
547     rextslog = 0
548     inprogress = 0
549     imax_pct = 25
550     icount = 64
551     ifree = 61
552     fdblocks = 62739235
553     frextents = 0
554     uquotino = 0
555     gquotino = 0
556     qflags = 0
557     flags = 0
558     shared_vn = 0
559     inoalignmt = 2
560     unit = 0
561     width = 0
562     dirblklog = 2
563     logsectlog = 0
564     logsectsize = 0
565     logsunit = 0
566     features2 = 8
567     ----
568
569
570     [[AG_Free_Space_Management]]
571     == AG Free Space Management
572
573     The XFS filesystem tracks free space in an allocation group using two B+trees.
574     One B+tree tracks space by block number, the second by the size of the free
575     space block. This scheme allows XFS to find quickly free space near a given
576     block or of a given size.
577
578     All block numbers, indexes, and counts are AG relative.
579
```

```

580  [[AG_Free_Space_Block]]
581  === AG Free Space Block
582
583  The second sector in an AG contains the information about the two free space
584  B+trees and associated free space information for the AG. The ``AG Free Space
585  Block'' also knows as the +AGF+, uses the following structure:
586
587  [source, c]
588  ----
589  struct xfs_agf {
590      __be32          agf_magicnum;
591      __be32          agf_versionnum;
592      __be32          agf_seqno;
593      __be32          agf_length;
594      __be32          agf_roots[XFS_BTNUM_AGF];
595      __be32          agf_levels[XFS_BTNUM_AGF];
596      __be32          agf_flfirst;
597      __be32          agf_fllast;
598      __be32          agf_flcount;
599      __be32          agf_freeblks;
600      __be32          agf_longest;
601      __be32          agf_btreeblks;
602
603      /* version 5 filesystem fields start here */
604      uuid_t          agf_uuid;
605      __be32          agf_rmap_blocks;
606      __be32          agf_refcount_blocks;
607      __be32          agf_refcount_root;
608      __be32          agf_refcount_level;
609      __be64          agf_spare64[14];
610
611      /* unlogged fields, written during buffer writeback. */
612      __be64          agf_lsn;
613      __be32          agf_crc;
614      __be32          agf_spare2;
615  };
616  ----
617
618  The rest of the bytes in the sector are zeroed. +XFS_BTNUM_AGF+ is set to 3:
619  index 0 for the free space B+tree indexed by block number; index 1 for the free
620  space B+tree indexed by extent size; and index 2 for the reverse-mapping
621  B+tree.
622
623  *agf_magicnum*::
624  Specifies the magic number for the AGF sector: ``XAGF'' (0x58414746).
625
626  *agf_versionnum*::
627  Set to +XFS_AGF_VERSION+ which is currently 1.
628

```

629 `*agf_seqno*::`
630 Specifies the AG number for the sector.
631
632 `*agf_length*::`
633 Specifies the size of the AG in filesystem blocks. For all AGs except the last,
634 this must be equal to the superblock's `+sb_agblocks+` value. For the last AG,
635 this could be less than the `+sb_agblocks+` value. It is this value that should
636 be used to determine the size of the AG.
637
638 `*agf_roots*::`
639 Specifies the block number for the root of the two free space B+trees and the
640 reverse-mapping B+tree, if enabled.
641
642 `*agf_levels*::`
643 Specifies the level or depth of the two free space B+trees and the
644 reverse-mapping B+tree, if enabled. For a fresh AG, this value will be one,
645 and the ```roots``` will point to a single leaf of level 0.
646
647 `*agf_flfirst*::`
648 Specifies the index of the first ```free list``` block. Free lists are covered in
649 more detail later on.
650
651 `*agf_fllast*::`
652 Specifies the index of the last ```free list``` block.
653
654 `*agf_flcount*::`
655 Specifies the number of blocks in the ```free list```.
656
657 `*agf_freeblks*::`
658 Specifies the current number of free blocks in the AG.
659
660 `*agf_longest*::`
661 Specifies the number of blocks of longest contiguous free space in the AG.
662
663 `*agf_btreeblks*::`
664 Specifies the number of blocks used for the free space B+trees. This is only
665 used if the `+XFS_SB_VERSION2_LAZYSBCOUNTBIT+` bit is set in `+sb_features2+`.
666
667 `*agf_uuid*::`
668 The UUID of this block, which must match either `+sb_uuid+` or `+sb_meta_uuid+`
669 depending on which features are set.
670
671 `*agf_rmap_blocks*::`
672 The size of the reverse mapping B+tree in this allocation group, in blocks.
673
674 `*agf_refcount_blocks*::`
675 The size of the reference count B+tree in this allocation group, in blocks.
676
677 `*agf_refcount_root*::`

```

678 Block number for the root of the reference count B+tree, if enabled.
679
680 *agf_refcount_level*::
681 Depth of the reference count B+tree, if enabled.
682
683 *agf_spare64*::
684 Empty space in the logged part of the AGF sector, for use for future features.
685
686 *agf_lsn*::
687 Log sequence number of the last AGF write.
688
689 *agf_crc*::
690 Checksum of the AGF sector.
691
692 *agf_spare2*::
693 Empty space in the unlogged part of the AGF sector.
694
695 [[AG_Free_Space_Btrees]]
696 === AG Free Space B+trees
697
698 The two Free Space B+trees store a sorted array of block offset and block
699 counts in the leaves of the B+tree. The first B+tree is sorted by the offset,
700 the second by the count or size.
701
702 Leaf nodes contain a sorted array of offset/count pairs which are also used for
703 node keys:
704
705 [source, c]
706 ----
707 struct xfs_alloc_rec {
708     __be32                ar_startblock;
709     __be32                ar_blockcount;
710 };
711 ----
712
713 *ar_startblock*::
714 AG block number of the start of the free space.
715
716 *ar_blockcount*::
717 Length of the free space.
718
719 Node pointers are an AG relative block pointer:
720
721 [source, c]
722 ----
723 typedef __be32 xfs_alloc_ptr_t;
724 ----
725
726 * As the free space tracking is AG relative, all the block numbers are only

```

```

727 32-bits.
728 * The +bb_magic+ value depends on the B+tree: ``ABTB'' (0x41425442) for the block
729 offset B+tree, ``ABTC'' (0x41425443) for the block count B+tree. On a v5
730 filesystem, these are ``AB3B'' (0x41423342) and ``AB3C'' (0x41423343),
731 respectively.
732 * The +xfs_btree_sblock_t+ header is used for intermediate B+tree node as well
733 as the leaves.
734 * For a typical 4KB filesystem block size, the offset for the +xfs_alloc_ptr_t+
735 array would be +0xab0+ (2736 decimal).
736 * There are a series of macros in +xfs_btree.h+ for deriving the offsets,
737 counts, maximums, etc for the B+trees used in XFS.
738
739 The following diagram shows a single level B+tree which consists of one leaf:
740
741 .Freespace B+tree with one leaf.
742 image::images/15a.png[]
743
744 With the intermediate nodes, the associated leaf pointers are stored in a
745 separate array about two thirds into the block. The following diagram
746 illustrates a 2-level B+tree for a free space B+tree:
747
748 .Multi-level freespace B+tree.
749 image::images/15b.png[]
750
751 [[AG_Free_List]]
752 == AG Free List
753
754 The AG Free List is located in the 4th sector of each AG and is known as the
755 AGFL. It is an array of AG relative block pointers for reserved space for
756 growing the free space B+trees. This space cannot be used for general user data
757 including inodes, data, directories and extended attributes.
758
759 With a freshly made filesystem, 4 blocks are reserved immediately after the free
760 space B+tree root blocks (blocks 4 to 7). As they are used up as the free space
761 fragments, additional blocks will be reserved from the AG and added to the free
762 list array. This size may increase as features are added.
763
764 As the free list array is located within a single sector, a typical device will
765 have space for 128 elements in the array (512 bytes per sector, 4 bytes per AG
766 relative block pointer). The actual size can be determined by using the
767 +XFS_AGFL_SIZE+ macro.
768
769 Active elements in the array are specified by the
770 xref:AG_Free_Space_Block[AGF's] +agf_flfirst+, +agf_fllast+ and +agf_flcount+
771 values. The array is managed as a circular list.
772
773 On a v5 filesystem, the following header precedes the free list entries:
774
775 [source, c]

```



```
776  ----
777  struct xfs_agfl {
778      __be32          agfl_magicnum;
779      __be32          agfl_seqno;
780      uuid_t          agfl_uuid;
781      __be64          agfl_lsn;
782      __be32          agfl_crc;
783  };
784  ----
785
786  *agfl_magicnum*::
787  Specifies the magic number for the AGFL sector: "XAFL" (0x5841464c).
788
789  *agfl_seqno*::
790  Specifies the AG number for the sector.
791
792  *agfl_uuid*::
793  The UUID of this block, which must match either +sb_uuid+ or +sb_meta_uuid+
794  depending on which features are set.
795
796  *agfl_lsn*::
797  Log sequence number of the last AGFL write.
798
799  *agfl_crc*::
800  Checksum of the AGFL sector.
801
802  On a v4 filesystem there is no header; the array of free block numbers begins
803  at the beginning of the sector.
804
805  .AG Free List layout
806  image::images/16.png[]
807
808  The presence of these reserved blocks guarantees that the free space B+trees
809  can be updated if any blocks are freed by extent changes in a full AG.
810
811  ==== xfs_db AGF Example
812
813  These examples are derived from an AG that has been deliberately fragmented.
814  The AGF:
815
816  ----
817  xfs_db> agf 0
818  xfs_db> p
819  magicnum = 0x58414746
820  versionnum = 1
821  seqno = 0
822  length = 3923122
823  bnoroot = 7
824  cntroot = 83343
```

```
825     bnolevel = 2
826     cntlevel = 2
827     flfirst = 22
828     fllast = 27
829     flcount = 6
830     freeblks = 3654234
831     longest = 3384327
832     btreeblks = 0
833     ----
834
835     In the AGFL, the active elements are from 22 to 27 inclusive which are obtained
836     from the +flfirst+ and +fllast+ values from the +agf+ in the previous example:
837
838     ----
839     xfs_db> agfl 0
840     xfs_db> p
841     bno[0-127] = 0:4 1:5 2:6 3:7 4:83342 5:83343 6:83344 7:83345 8:83346 9:83347
842                  10:4 11:5 12:80205 13:80780 14:81496 15:81766 16:83346 17:4 18:5
843                  19:80205 20:82449 21:81496 22:81766 23:82455 24:80780 25:5
844                  26:80205 27:83344
845     ----
846
847     The root block of the free space B+tree sorted by block offset is found in the
848     AGF's +bnoroot+ value:
849
850     ----
851     xfs_db> fsblock 7
852     xfs_db> type bnobt
853     xfs_db> p
854     magic = 0x41425442
855     level = 1
856     numrecs = 4
857     leftsib = null
858     rightsib = null
859     keys[1-4] = [startblock,blockcount]
860                1:[12,16] 2:[184586,3] 3:[225579,1] 4:[511629,1]
861     ptrs[1-4] = 1:2 2:83347 3:6 4:4
862     ----
863
864     Blocks 2, 83347, 6 and 4 contain the leaves for the free space B+tree by
865     starting block. Block 2 would contain offsets 12 up to but not including 184586
866     while block 4 would have all offsets from 511629 to the end of the AG.
867
868     The root block of the free space B+tree sorted by block count is found in the
869     AGF's +cntroot+ value:
870
871     ----
872     xfs_db> fsblock 83343
873     xfs_db> type cntbt
```

```
874 xfs_db> p
875 magic = 0x41425443
876 level = 1
877 numrecs = 4
878 leftsib = null
879 rightsib = null
880 keys[1-4] = [blockcount,startblock]
881           1:[1,81496] 2:[1,511729] 3:[3,191875] 4:[6,184595]
882 ptrs[1-4] = 1:3 2:83345 3:83342 4:83346
883 ----
884
885 The leaf in block 3, in this example, would only contain single block counts.
886 The offsets are sorted in ascending order if the block count is the same.
887
888 Inspecting the leaf in block 83346, we can see the largest block at the end:
889
890 ----
891 xfs_db> fsblock 83346
892 xfs_db> type cntbt
893 xfs_db> p
894 magic = 0x41425443
895 level = 0
896 numrecs = 344
897 leftsib = 83342
898 rightsib = null
899 recs[1-344] = [startblock,blockcount]
900             1:[184595,6] 2:[187573,6] 3:[187776,6]
901             ...
902             342:[513712,755] 343:[230317,258229] 344:[538795,3384327]
903 ----
904
905 The longest block count (3384327) must be the same as the AGF's +longest+ value.
906
907 [[AG_Inode_Management]]
908 == AG Inode Management
909
910 [[Inode_Numbers]]
911 == Inode Numbers
912
913 Inode numbers in XFS come in two forms: AG relative and absolute.
914
915 AG relative inode numbers always fit within 32 bits. The number of bits actually
916 used is determined by the sum of the xref:Superblocks[superblock's] +sb_inoplog+
917 and +sb_agblklog+ values. Relative inode numbers are found within the AG's inode
918 structures.
919
920 Absolute inode numbers include the AG number in the high bits, above the bits
921 used for the AG relative inode number. Absolute inode numbers are found in
922 xref:Directories[directory] entries and the superblock.
```

```
923
924 .Inode number formats
925 image::images/18.png[]
926
927 [[Inode_Information]]
928 === Inode Information
929
930 Each AG manages its own inodes. The third sector in the AG contains information
931 about the AG's inodes and is known as the AGI.
932
933 The AGI uses the following structure:
934
935 [source, c]
936 ----
937 struct xfs_agi {
938     __be32          agi_magicnum;
939     __be32          agi_versionnum;
940     __be32          agi_seqno
941     __be32          agi_length;
942     __be32          agi_count;
943     __be32          agi_root;
944     __be32          agi_level;
945     __be32          agi_freecount;
946     __be32          agi_newino;
947     __be32          agi_dirino;
948     __be32          agi_unlinked[64];
949
950     /*
951      * v5 filesystem fields start here; this marks the end of logging region 1
952      * and start of logging region 2.
953      */
954     uuid_t          agi_uuid;
955     __be32          agi_crc;
956     __be32          agi_pad32;
957     __be64          agi_lsn;
958
959     __be32          agi_free_root;
960     __be32          agi_free_level;
961
962     __be32          agi_iblocks;
963     __be32          agi_fblocks;
964
965 }
966 ----
967 *agi_magicnum*::
968 Specifies the magic number for the AGI sector: ``XAGI'' (0x58414749).
969
970 *agi_versionnum*::
971 Set to +XFS_AGI_VERSION+ which is currently 1.
```

```
972
973 *agi_seqno*::
974 Specifies the AG number for the sector.
975
976 *agi_length*::
977 Specifies the size of the AG in filesystem blocks.
978
979 *agi_count*::
980 Specifies the number of inodes allocated for the AG.
981
982 *agi_root*::
983 Specifies the block number in the AG containing the root of the inode B+tree.
984
985 *agi_level*::
986 Specifies the number of levels in the inode B+tree.
987
988 *agi_freecount*::
989 Specifies the number of free inodes in the AG.
990
991 *agi_newino*::
992 Specifies AG-relative inode number of the most recently allocated chunk.
993
994 *agi_dirino*::
995 Deprecated and not used, this is always set to NULL (-1).
996
997 *agi_unlinked[64]*::
998 Hash table of unlinked (deleted) inodes that are still being referenced. Refer
999 to xref:Unlinked_Pointer[unlinked list pointers] for more information.
1000
1001 *agi_uuid*::
1002 The UUID of this block, which must match either +sb_uuid+ or +sb_meta_uuid+
1003 depending on which features are set.
1004
1005 *agi_crc*::
1006 Checksum of the AGI sector.
1007
1008 *agi_pad32*::
1009 Padding field, otherwise unused.
1010
1011 *agi_lsn*::
1012 Log sequence number of the last write to this block.
1013
1014 *agi_free_root*::
1015 Specifies the block number in the AG containing the root of the free inode
1016 B+tree.
1017
1018 *agi_free_level*::
1019 Specifies the number of levels in the free inode B+tree.
1020
```

```
1021 *agi_iblocks*::
1022 The number of blocks in the inode B+tree, including the root.
1023 This field is zero if the +XFS_SB_FEAT_RO_COMPAT_INOBT CNT+ feature is not
1024 enabled.
1025
1026 *agi_fblocks*::
1027 The number of blocks in the free inode B+tree, including the root.
1028 This field is zero if the +XFS_SB_FEAT_RO_COMPAT_INOBT CNT+ feature is not
1029 enabled.
1030
1031 [[Inode_Btrees]]
1032 == Inode B+trees
1033
1034 Inodes are traditionally allocated in chunks of 64, and a B+tree is used to
1035 track these chunks of inodes as they are allocated and freed. The block
1036 containing root of the B+tree is defined by the AGI's +agi_root+ value. If the
1037 +XFS_SB_FEAT_RO_COMPAT_FINOBT+ feature is enabled, a second B+tree is used to
1038 track the chunks containing free inodes; this is an optimization to speed up
1039 inode allocation.
1040
1041 The B+tree header for the nodes and leaves use the +xfs_btree_sblock+ structure
1042 which is the same as the header used in the xref:AG_Free_Space_Btrees[AGF
1043 B+trees].
1044
1045 The magic number of the inode B+tree is ``IABT'' (0x49414254). On a v5
1046 filesystem, the magic number is ``IAB3'' (0x49414233).
1047
1048 The magic number of the free inode B+tree is ``FIBT'' (0x46494254). On a v5
1049 filesystem, the magic number is ``FIB3'' (0x46494254).
1050
1051 Leaves contain an array of the following structure:
1052
1053 [source,c]
1054 ----
1055 struct xfs_inobt_rec {
1056     __be32                ir_startino;
1057     __be32                ir_freecount;
1058     __be64                ir_free;
1059 };
1060 ----
1061
1062 *ir_startino*::
1063 The lowest-numbered inode in this chunk.
1064
1065 *ir_freecount*::
1066 Number of free inodes in this chunk.
1067
1068 *ir_free*::
1069 A 64 element bitmap showing which inodes in this chunk are free.
```

```
1070
1071 Nodes contain key/pointer pairs using the following types:
1072
1073 [source,c]
1074 ----
1075 struct xfs_inobt_key {
1076     __be32                ir_startino;
1077 };
1078 typedef __be32 xfs_inobt_ptr_t;
1079 ----
```

```
1080
1081 The following diagram illustrates a single level inode B+tree:
1082
```

```
1083 .Single Level inode B+tree
1084 image::images/20a.png[]
```

```
1085
```

```
1086
```

```
1087 And a 2-level inode B+tree:
```

```
1088
```

```
1089 .Multi-Level inode B+tree
1090 image::images/20b.png[]
```

```
1091
```

```
1092
```

```
1093 === xfs_db AGI Example
```

```
1094
```

```
1095 This is an AGI of a freshly populated filesystem:
```

```
1096
```

```
1097 ----
```

```
1098 xfs_db> agi 0
```

```
1099 xfs_db> p
```

```
1100 magicnum = 0x58414749
```

```
1101 versionnum = 1
```

```
1102 seqno = 0
```

```
1103 length = 825457
```

```
1104 count = 5440
```

```
1105 root = 3
```

```
1106 level = 1
```

```
1107 freecount = 9
```

```
1108 newino = 5792
```

```
1109 dirino = null
```

```
1110 unlinked[0-63] =
```

```
1111 uuid = 3dfa1e5c-5a5f-4ca2-829a-000e453600fe
```

```
1112 lsn = 0x1000032c2
```

```
1113 crc = 0x14cb7e5c (correct)
```

```
1114 free_root = 4
```

```
1115 free_level = 1
```

```
1116 ----
```

```
1117
```

```
1118 From this example, we see that the inode B+tree is rooted at AG block 3 and
```

```
1119 that the free inode B+tree is rooted at AG block 4. Let's look at the
1120 inode B+tree:
1121
1122 ----
1123 xfs_db> addr root
1124 xfs_db> p
1125 magic = 0x49414233
1126 level = 0
1127 numrecs = 85
1128 leftsib = null
1129 rightsib = null
1130 bno = 24
1131 lsn = 0x1000032c2
1132 uuid = 3dfa1e5c-5a5f-4ca2-829a-000e453600fe
1133 owner = 0
1134 crc = 0x768f9592 (correct)
1135 recs[1-85] = [startino, freecount, free]
1136      1:[96,0,0] 2:[160,0,0] 3:[224,0,0] 4:[288,0,0]
1137      5:[352,0,0] 6:[416,0,0] 7:[480,0,0] 8:[544,0,0]
1138      9:[608,0,0] 10:[672,0,0] 11:[736,0,0] 12:[800,0,0]
1139      ...
1140      85:[5792,9,0xff80000000000000]
1141 ----
1142
1143 Most of the inode chunks on this filesystem are totally full, since the +free+
1144 value is zero. This means that we ought to expect inode 160 to be linked
1145 somewhere in the directory structure. However, notice that 0xff80000000000000
1146 in record 85 -- this means that we would expect inode 5847 to be free. Moving
1147 on to the free inode B+tree, we see that this is indeed the case:
1148
1149 ----
1150 xfs_db> addr free_root
1151 xfs_db> p
1152 magic = 0x46494233
1153 level = 0
1154 numrecs = 1
1155 leftsib = null
1156 rightsib = null
1157 bno = 32
1158 lsn = 0x1000032c2
1159 uuid = 3dfa1e5c-5a5f-4ca2-829a-000e453600fe
1160 owner = 0
1161 crc = 0x338af88a (correct)
1162 recs[1] = [startino, freecount, free] 1:[5792,9,0xff80000000000000]
1163 ----
1164
1165 Observe also that the AGI's +agi_newino+ points to this chunk, which has never
1166 been fully allocated.
1167
```



```
1168 [[Sparse_Inodes]]
1169 == Sparse Inodes
1170
1171 As mentioned in the previous section, XFS allocates inodes in chunks of 64. If
1172 there are no free extents large enough to hold a full chunk of 64 inodes, the
1173 inode allocation fails and XFS claims to have run out of space. On a
1174 filesystem with highly fragmented free space, this can lead to out of space
1175 errors long before the filesystem runs out of free blocks.
1176
1177 The sparse inode feature tracks inode chunks in the inode B+tree as if they
1178 were full chunks but uses some previously unused bits in the freecount field to
1179 track which parts of the inode chunk are not allocated for use as inodes. This
1180 allows XFS to allocate inodes one block at a time if absolutely necessary.
1181
1182 The inode and free inode B+trees operate in the same manner as they do without
1183 the sparse inode feature; the B+tree header for the nodes and leaves use the
1184 +xfs_btree_sblock+ structure which is the same as the header used in the
1185 xref:AG_Free_Space_Btrees[AGF B+trees].
1186
1187 It is theoretically possible for a sparse inode B+tree record to reference
1188 multiple non-contiguous inode chunks.
1189
1190 Leaves contain an array of the following structure:
1191
1192 [source,c]
1193 ----
1194 struct xfs_inobt_rec {
1195     __be32          ir_startino;
1196     __be16          ir_holemask;
1197     __u8            ir_count;
1198     __u8            ir_freecount;
1199     __be64          ir_free;
1200 };
1201 ----
1202
1203 *ir_startino*::
1204 The lowest-numbered inode in this chunk, rounded down to the nearest multiple
1205 of 64, even if the start of this chunk is sparse.
1206
1207 *ir_holemask*::
1208 A 16 element bitmap showing which parts of the chunk are not allocated to
1209 inodes. Each bit represents four inodes; if a bit is marked here, the
1210 corresponding bits in ir_free must also be marked.
1211
1212 *ir_count*::
1213 Number of inodes allocated to this chunk.
1214
1215 *ir_freecount*::
1216 Number of free inodes in this chunk.
```

```
1217
1218  *ir_free*::
1219  A 64 element bitmap showing which inodes in this chunk are not available for
1220  allocation.
1221
1222  === xfs_db Sparse Inode AGI Example
1223
1224  This example derives from an AG that has been deliberately fragmented.  The
1225  inode B+tree:
1226
1227  ----
1228  xfs_db> agi 0
1229  xfs_db> p
1230  magicnum = 0x58414749
1231  versionnum = 1
1232  seqno = 0
1233  length = 6400
1234  count = 10432
1235  root = 2381
1236  level = 2
1237  freecount = 0
1238  newino = 14912
1239  dirino = null
1240  unlinked[0-63] =
1241  uuid = b9b4623b-f678-4d48-8ce7-ce08950e3cd6
1242  lsn = 0x600000ac4
1243  crc = 0xef550dbc (correct)
1244  free_root = 4
1245  free_level = 1
1246  ----
1247
1248  This AGI was formatted on a v5 filesystem; notice the extra v5 fields.  So far
1249  everything else looks much the same as always.
1250
1251  ----
1252  xfs_db> addr root
1253  magic = 0x49414233
1254  level = 1
1255  numrecs = 2
1256  leftsib = null
1257  rightsib = null
1258  bno = 19048
1259  lsn = 0x50000192b
1260  uuid = b9b4623b-f678-4d48-8ce7-ce08950e3cd6
1261  owner = 0
1262  crc = 0xd98cd2ca (correct)
1263  keys[1-2] = [startino] 1:[128] 2:[35136]
1264  ptrs[1-2] = 1:3 2:2380
1265  xfs_db> addr ptrs[1]
```

```

1266 xfs_db> p
1267 magic = 0x49414233
1268 level = 0
1269 numrecs = 159
1270 leftsib = null
1271 rightsib = 2380
1272 bno = 24
1273 lsn = 0x600000ac4
1274 uuid = b9b4623b-f678-4d48-8ce7-ce08950e3cd6
1275 owner = 0
1276 crc = 0x836768a6 (correct)
1277 recs[1-159] = [startino,holemask,count,freecount,free]
1278     1:[128,0,64,0,0]
1279     2:[14912,0xff,32,0,0xffffffff]
1280     3:[15040,0,64,0,0]
1281     4:[15168,0xff00,32,0,0xffffffff00000000]
1282     5:[15296,0,64,0,0]
1283     6:[15424,0xff,32,0,0xffffffff]
1284     7:[15552,0,64,0,0]
1285     8:[15680,0xff00,32,0,0xffffffff00000000]
1286     9:[15808,0,64,0,0]
1287    10:[15936,0xff,32,0,0xffffffff]
1288 ----
1289
1290 Here we see the difference in the inode B+tree records. For example, in record
1291 2, we see that the holemask has a value of 0xff. This means that the first
1292 sixteen inodes in this chunk record do not actually map to inode blocks; the
1293 first inode in this chunk is actually inode 14944:
1294
1295 ----
1296 xfs_db> inode 14912
1297 Metadata corruption detected at block 0x3a40/0x2000
1298 ...
1299 Metadata CRC error detected for ino 14912
1300 xfs_db> p core.magic
1301 core.magic = 0
1302 xfs_db> inode 14944
1303 xfs_db> p core.magic
1304 core.magic = 0x494e
1305 ----
1306
1307 The chunk record also indicates that this chunk has 32 inodes, and that the
1308 missing inodes are also ``free''.
1309
1310 [[Real-time_Devices]]
1311 == Real-time Devices
1312
1313 The performance of the standard XFS allocator varies depending on the internal
1314 state of the various metadata indices enabled on the filesystem. For

```

1315 applications which need to minimize the jitter of allocation latency, XFS
1316 supports the notion of a ``real-time device''. This is a special device
1317 separate from the regular filesystem where extent allocations are tracked with
1318 a bitmap and free space is indexed with a two-dimensional array. If an inode
1319 is flagged with +XFS_DIFLAG_REALTIME+, its data will live on the real time
1320 device. The metadata for real time devices is discussed in the section about
1321 xref:Real-time_Inodes[real time inodes].
1322
1323 By placing the real time device (and the journal) on separate high-performance
1324 storage devices, it is possible to reduce most of the unpredictability in I/O
1325 response times that come from metadata operations.
1326
1327 None of the XFS per-AG B+trees are involved with real time files. It is not
1328 possible for real time files to share data blocks.