

XFS uses several different directory structures depending on the size of the directory. For testing purposes, I created three directories— one with 5 files, one with 50, and one with 5000 file entries. Small directories have their data stored in the inode. In this installment we'll examine the inode of the directory that contains only five files.

0x000	49	4E	41	ED	03	01	00	00	00	00	00	00	00	00	00	00	I NAí.....
0x010	00	00	00	02	00	00	00	00	00	00	00	00	00	00	00	00
0x020	5A	EF	84	6C	11	16	6A	BB	5A	EF	84	67	03	38	E8	90	Zï.l..j»Zï.g.8è.
0x030	5A	EF	84	67	03	38	E8	90	00	00	00	00	00	00	00	66	Zï.g.8è.....f
0x040	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0x050	00	00	23	01	00	00	00	00	00	00	00	00	5A	52	79	A1	..#.....ZRy
0x060	FF	FF	FF	FF	27	FA	73	4A	00	00	00	00	00	00	00	01	ÿÿÿÿ'úsJ.....
0x070	00	00	00	20	00	00	7F	FC	00	00	00	00	00	00	00	00ü.....
0x080	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0x090	5A	EF	84	47	13	FF	46	D3	00	00	00	00	04	08	E5	96	Zï.G.ÿFÓ.....å.
0x0A0	E5	6C	3B	41	CA	03	4B	41	B1	5C	DD	60	9C	B7	DA	71	ål;AÊ.KA±\Ý`..Úq
0x0B0	05	00	04	15	9F	A1	0C	00	60	30	31	5F	73	6D	61	6Ci .`01_sma
0x0C0	6C	66	69	6C	65	01	04	17	97	9D	0A	00	78	30	32	5F	lfile..... x02_
0x0D0	62	69	67	66	69	6C	65	01	04	17	97	9E	0C	00	90	30	bigfile..... .0
0x0E0	33	5F	73	6D	61	6C	6C	66	69	6C	65	01	04	17	97	9F	3_smallfile.....
0x0F0	0A	00	A8	30	34	5F	62	69	67	66	69	6C	65	01	04	17	.`04_bigfile...
0x100	A1	54	0C	00	C0	30	35	5F	73	6D	61	6C	6C	66	69	6C	iT .A05_smallfil
0x110	65	01	04	17	A1	55	00	00	00	00	00	00	00	00	00	00	e...iU.....
0x120	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0x130	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0x140	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0x150	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0x160	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0x170	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0x180	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0x190	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0x1A0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0x1B0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
0x1C0	00	00	00	00	00	00	00	00	00	34	01	00	07	26	04	734...&.s
0x1D0	65	6C	69	6E	75	78	75	6E	63	6F	6E	66	69	6E	65	64	elinuxunconfined
0x1E0	5F	75	3A	6F	62	6A	65	63	74	5F	72	3A	61	64	6D	69	_u:object_r:admi
0x1F0	6E	5F	68	6F	6D	65	5F	74	3A	73	30	00	00	73	30	00	n_home_t:s0..s0.

 $\frac{1}{4}$

Resident directory data is stored as a “short form” directory structure starting at byte offset 176, right after the inode core. First we have a brief header:

176	Number of directory entries	5
177	Number of dir entries needing 64-bit inodes	0
178-181	Inode of parent	0x04159fa1

First we have a byte tracking the number of directory entries to follow the header. The next byte tracks how many directory entries require 64 bits for inode data. As we saw in [Part 1](https://righteousit.wordpress.com/2018/05/21/xfs-part-1-superblock/) (<https://righteousit.wordpress.com/2018/05/21/xfs-part-1-superblock/>) of this series, XFS uses variable length addresses for blocks and inodes. In our file system, we need less than 32 bits to store these addresses, so there are no directory entries requiring 64-bit inodes. This means the directory data will use 32 bits to store inodes in order to save space.

This has an immediate impact because the next entry in the header is the inode of the parent directory. Since byte 177 is zero, this field will be 32 bits. If byte 177 was non-zero, then all inode entries in the header and directory entries would be 64-bit.

The parent inode field in the header is the equivalent of the usual “.” link in the directory. The current directory inode (the “.” link) is found in the inode core in bytes 152-159. The short form directory simply uses these values and does not have explicit “.” and “..” entries.

After the header come a series of variable length directory entries, packed as tightly as possible with no alignment constraints. Entries are added to the directory in order of file creation and are not sorted in any way.

Here is a description of the fields and a breakdown of the values in the five directories in this inode:

Len (Bytes)	Field
1	Length of file name (in bytes)
2	Entry offset in non short form directory
varies	Characters in file name
1	File type
4 or 8	Absolute inode address

Len	Offset	Name	Type	Inode
===	=====	====	=====	=====
12	0x0060	01_smallfile	01	0x0417979d
10	0x0078	02_bigfile	01	0x0417979e
12	0x0090	03_smallfile	01	0x0417979f
10	0x00a8	04_bigfile	01	0x0417a154
12	0x00c0	05_smallfile	01	0x0417a155

First we have a single byte for the file name length in bytes. Like other Unix file systems, there is a 255 character file name limit.

The next two bytes are based on the byte offset the directory entry would have if it were a normal XFS directory entry and not packed into a short form directory in the inode. In a normal directory block, directory entries are 64-bit aligned and start at byte offset 96 (0x60) following the directory header and “.” and “..” entries. The directory entries here are all 18 or 20 bytes long, which means they would consume 24 bytes (0x18) in a normal directory block. Using a consistent numbering scheme for the offset makes it easier to write code that iterates through directory entries, even though the offsets don’t match the actual offset of each directory entry in the short form style.

Next we have the characters in the file name followed by a single byte for the file type. The file type is included in the directory entry so that commands like “ls -F” don’t have to open each inode to get the file type information. The file type values in the directory entry do not use the same number scheme as the file type in the inode. Here are the expected values for directory entries:

- 1 Regular file
- 2 Directory
- 3 Character special device
- 4 Block special device
- 5 FIFO
- 6 Socket
- 7 Symlink

Finally there is a field to hold the inode associated with the file name. In our example, these inode entries are 32 bits. 64-bit inode fields will be used if the directory header indicates they are needed.

Deleting a File

When a file is deleted from (or added to) a directory, the mtime and ctime in the directory’s inode core are updated. The directory file size changes (bytes 56-63). The CRC32 checksum and the logfile sequence number fields are updated.

In the data fork, all directory entries after the deleted entry are shifted downwards, completely overwriting the deleted entry. Here’s what the directory entries look like after “03_smallfile”– the third entry in the original directory– is deleted:

0x0B0	04	00	04	15	9F	A1	0C	00	60	30	31	5F	73	6D	61	6Ci .`01_sma
0x0C0	6C	66	69	6C	65	01	04	17	97	9D	0A	00	78	30	32	5F	lfile..... .x02_
0x0D0	62	69	67	66	69	6C	65	01	04	17	97	9E	0A	00	A8	30	bigfile..... .`0
0x0E0	34	5F	62	69	67	66	69	6C	65	01	04	17	A1	54	0C	00	4_bigfile...iT .
0x0F0	C0	30	35	5F	73	6D	61	6C	6C	66	69	6C	65	01	04	17	A05_smallfile...
0x100	A1	55	0C	00	C0	30	35	5F	73	6D	61	6C	6C	66	69	6C	iU .A05_smallfil
0x110	65	01	04	17	A1	55	00	00	00	00	00	00	00	00	00	00	e...iU.....

The four remaining directory entries are highlighted above. However, after those entries you can clearly see the residue of the entry for “05_smallfile” from the original directory. So as short-form directories shrink, they leave behind entries in the unused “inode slack”. In this case the residue is for a file entry that still exists in the directory, but it’s possible that we might get residue of entries deleted from the end of the directory list.

When Directories Grow Up

Another place you can see short form directory residue is when the directory gets large enough that it needs to move out to blocks on disk. I created a sample directory that initially had five files and confirmed that it was being stored as a short form directory in the inode. Then I added 45 more files to the directory, which made a

[illegible]

The format of directories changes significantly when directory entries move out into disk blocks. In our next installment we will examine the structures in these larger directories.

Posted in [Forensics](#) Tagged [File Systems](#), [Linux](#), [XFS](#)



Published by Hal Pomeranz

Create a free website or blog at WordPress.com.